

# Final Exam Review

Fall 2017

# Outline

- Chapter 0 --- Basic Linux Commands and C Programming
- Chapter 1 --- Operating System Overview
- Chapter 2 --- Linux and Process Overview
- Chapter 3 --- Processes concepts and implementation
- Chapter 4 --- Threads and System Calls
- Chapter 5 --- Concurrency
- Chapter 6 --- Deadlock
- Chapter 7 --- Memory Management
- Chapter 8 --- Virtual Memory and Virtual Page Allocation

# Chapter 0

Basic Linux Commands and C Programming

# Basic Linux Commands

- cat
- ls
- gcc
- make
- chmod
- scp
- tar
- ssh
- scp

# Universality of I/O

**\$ ./copy test test.old**

*Copy a regular file*

**\$ ./copy a.txt /dev/tty**

*Copy a regular file to this terminal*

**\$ ./copy /dev/tty b.txt**

*Copy input from this terminal to a regular file*

**\$ ./copy /dev/pts/16 /dev/tty**

*Copy input from another terminal*

# FILES in C

- `#include <stdio.h>`
- `FILE *`: pointer to a file object
  - The `FILE` object is a special data structure that keeps track of the current file information
- Special files
  - **`stdin`**: Standard input
  - **`stdout`**: standard output
  - **`stderr`**: Standard error

File descriptor	Purpose	POSIX name	<i>stdio</i> stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

# FILES in C

- `#include <stdio.h>`
- `FILE *`: pointer to a file object
  - The `FILE` object is a special data structure that keeps track of the current file information
- Special files
  - **`stdin`**: Standard input
  - **`stdout`**: standard output
  - **`stderr`**: Standard error

File descriptor	Purpose	POSIX name	<i>stdio</i> stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

# Chapter 1

Operating Systems Overview




# Basic Elements

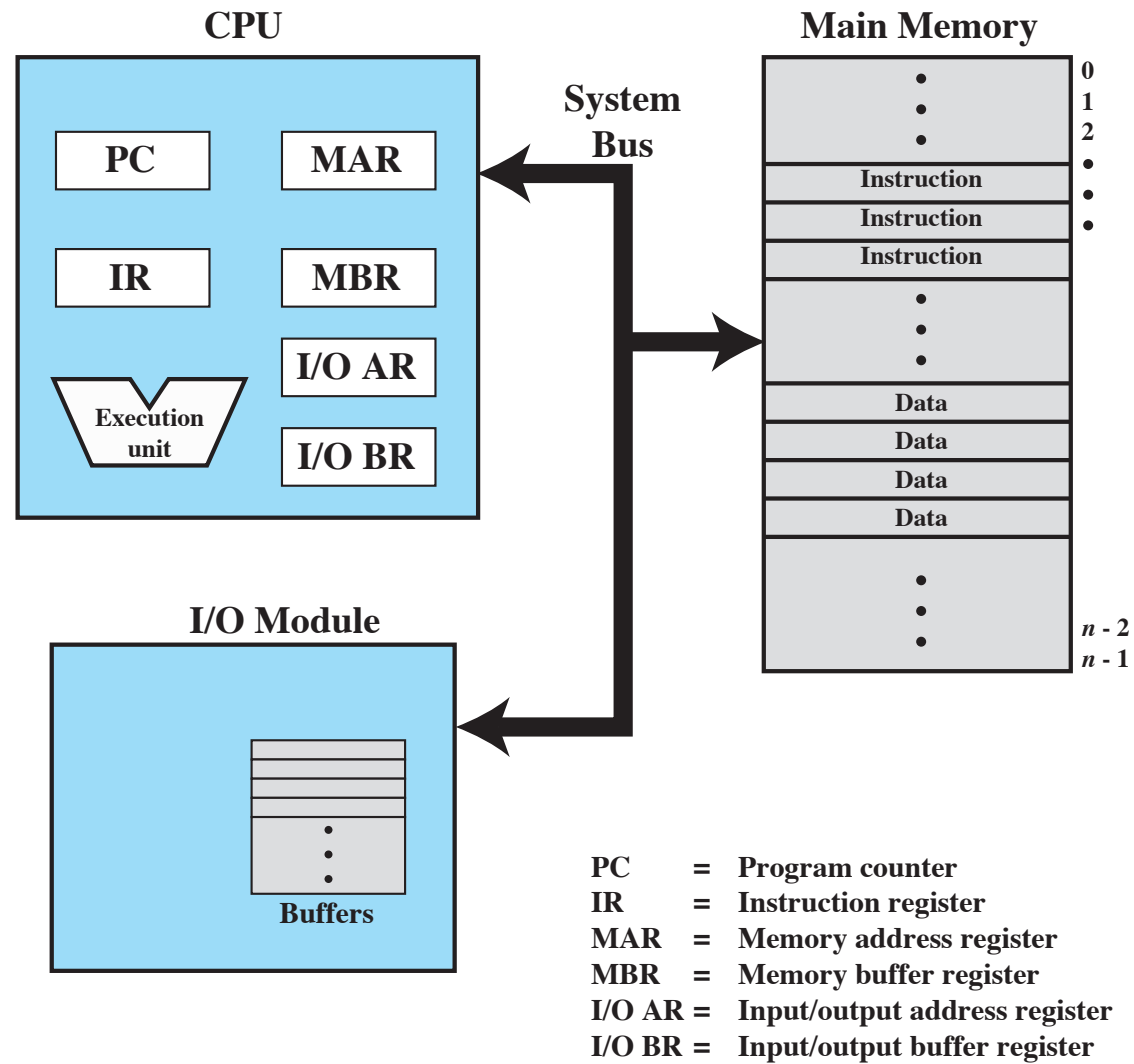
**Processor**

**I/O  
Modules**

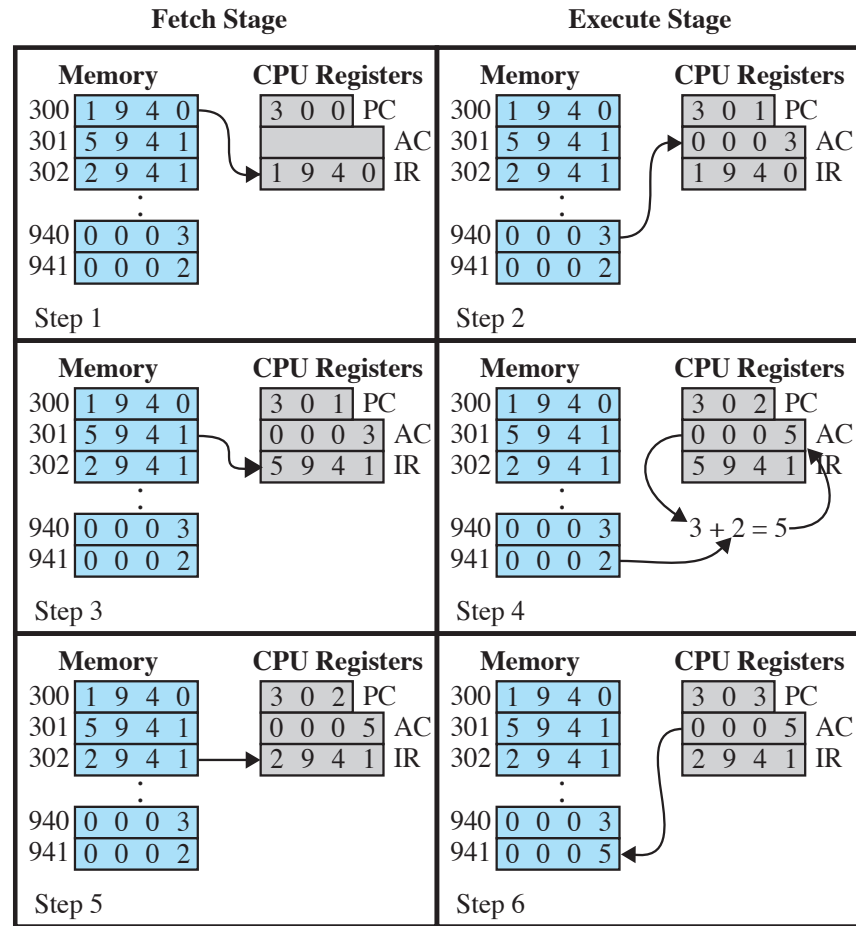
**Main  
Memory**

**System  
Bus**





**Figure 1.1 Computer Components: Top-Level View**



**Figure 1.4 Example of Program Execution**  
 (contents of memory and registers in hexadecimal)

# Interrupts

- Interrupt the normal sequencing of the processor
- Provided to improve processor utilization
  - most I/O devices are slower than the processor
  - processor must pause to wait for device
  - wasteful use of the processor



## Table 1.1    Classes of Interrupts

**Program**                      Generated by some condition that occurs as a result of an instruction execution, such as arithmetic overflow, division by zero, attempt to execute an illegal machine instruction, and reference outside a user's allowed memory space.

**Timer**                        Generated by a timer within the processor. This allows the operating system to perform certain functions on a regular basis.

**I/O**                            Generated by an I/O controller, to signal normal completion of an operation or to signal a variety of error conditions.

**Hardware failure**            Generated by a failure, such as power failure or memory parity error.

# The Memory Hierarchy

- Going down the hierarchy:
  - decreasing cost per bit
  - increasing capacity
  - increasing access time
  - decreasing frequency of access to the memory by the processor




Figure 1.14 The Memory Hierarchy

# Principle of Locality


- Memory references by the processor tend to cluster
- Data is organized so that the percentage of accesses to each successively lower level is substantially less than that of the level above
- Can be applied across more than two levels of memory

# Chapter 2

Linux and Process Overview



# Process



- Fundamental to the structure of operating systems

A *process* can be defined as:

a program in execution

an instance of a running program

the entity that can be assigned to, and executed on, a processor

a unit of activity characterized by a single sequential thread of execution, a current state, and an associated set of system resources

# LINUX Overview

- Started out as a UNIX variant for the IBM PC
- Linus Torvalds, a Finnish student of computer science, wrote the initial version
- Linux was first posted on the Internet in 1991
- Today it is a full-featured UNIX system that runs on several platforms
- Is free and the source code is available
- Key to success has been the availability of free software packages
- Highly modular and easily configured

# Modular Monolithic Kernel

- Includes virtually all of the OS functionality in one large block of code that runs as a single process with a single address space
- All the functional components of the kernel have access to all of its internal data structures and routines
- Linux is structured as a collection of modules

## Loadable Modules

- Relatively independent blocks
- A module is an object file whose code can be linked to and unlinked from the kernel at runtime
- A module is executed in kernel mode on behalf of the current process
- Have two important characteristics:
  - dynamic linking
  - stackable modules

# Chapter 3

Process Concepts and Implementation

# Process Elements

- Two essential elements of a process are:

## Program code


- which may be shared with other processes that are executing the same program

## A set of data associated with that code

- when the processor begins to execute the program code, we refer to this executing entity as a *process*

# Process Elements

- While the program is executing, this process can be uniquely characterized by a number of elements, including:



# Five-State Process Model

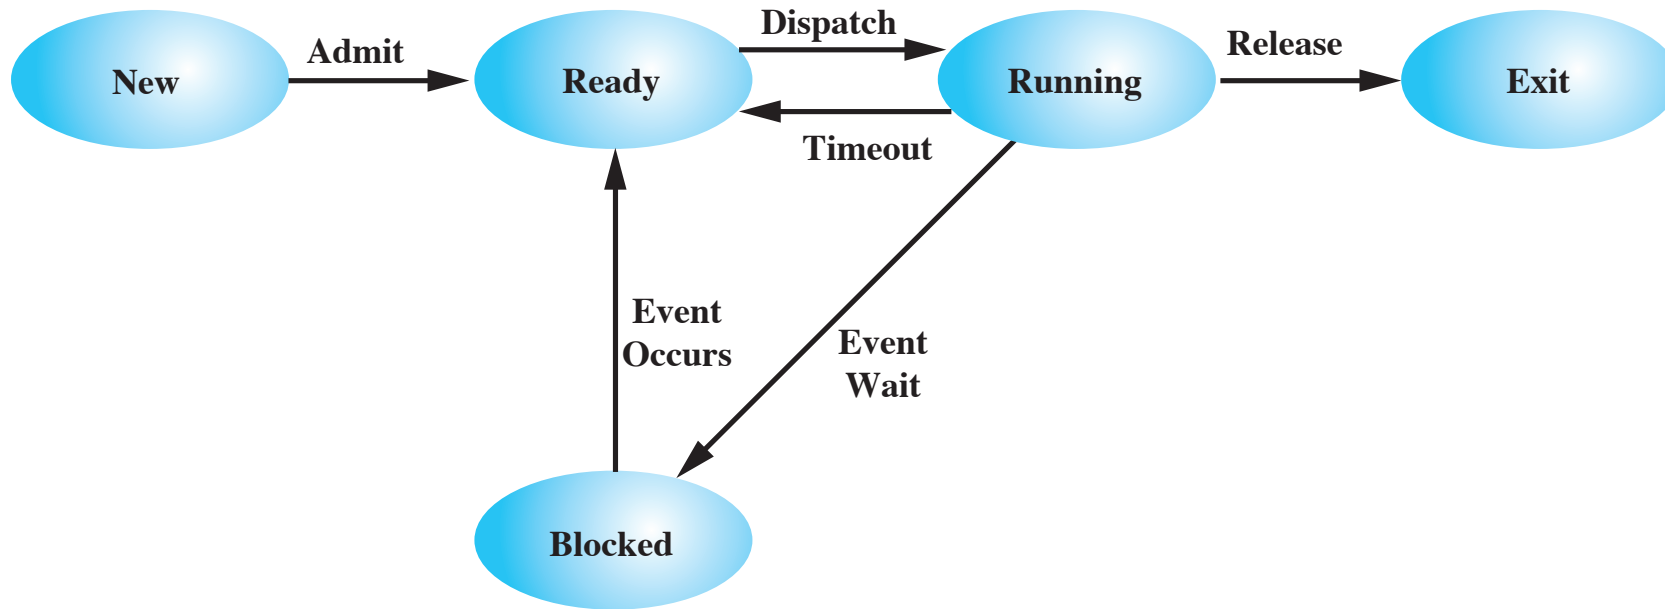

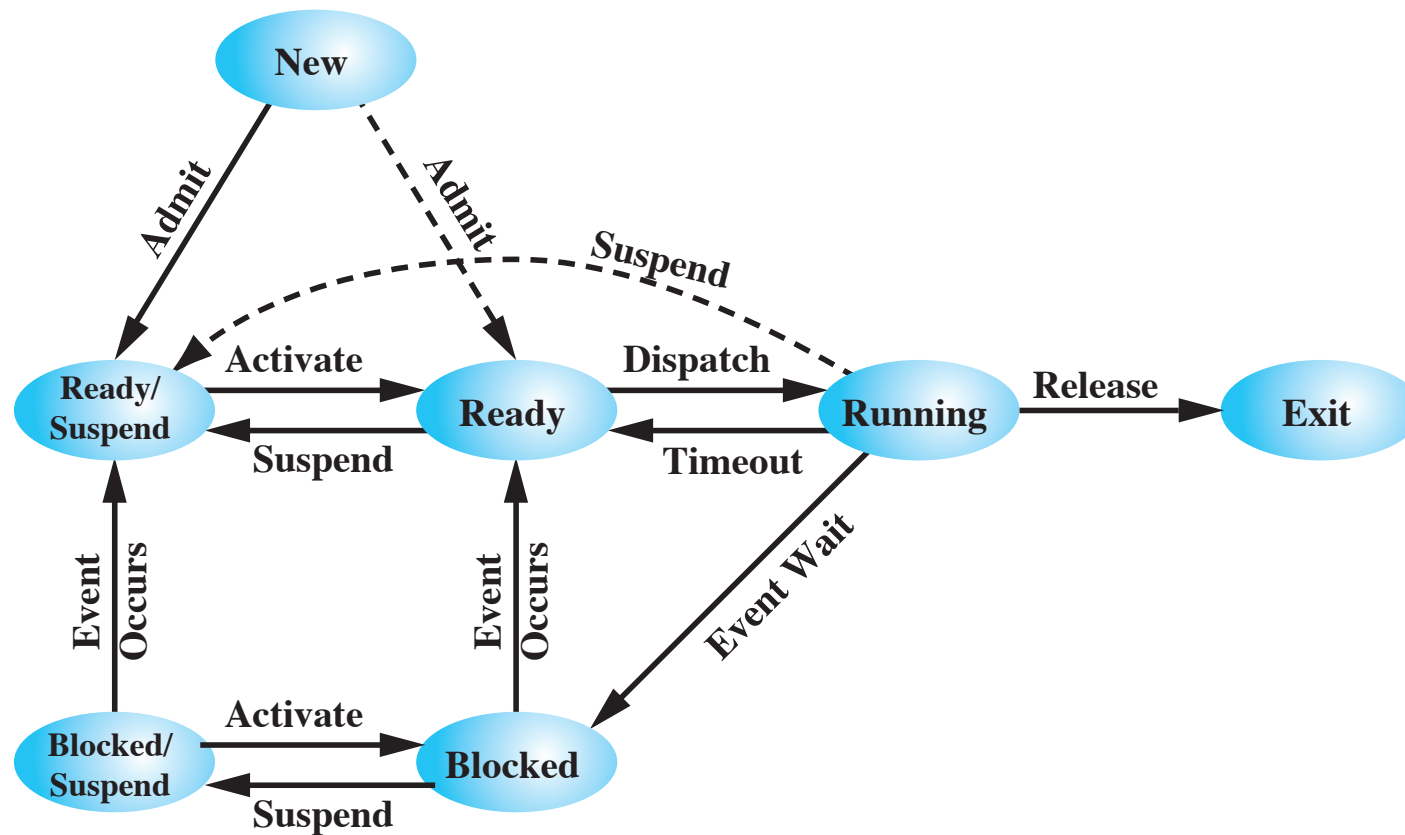


Figure 3.6 Five-State Process Model



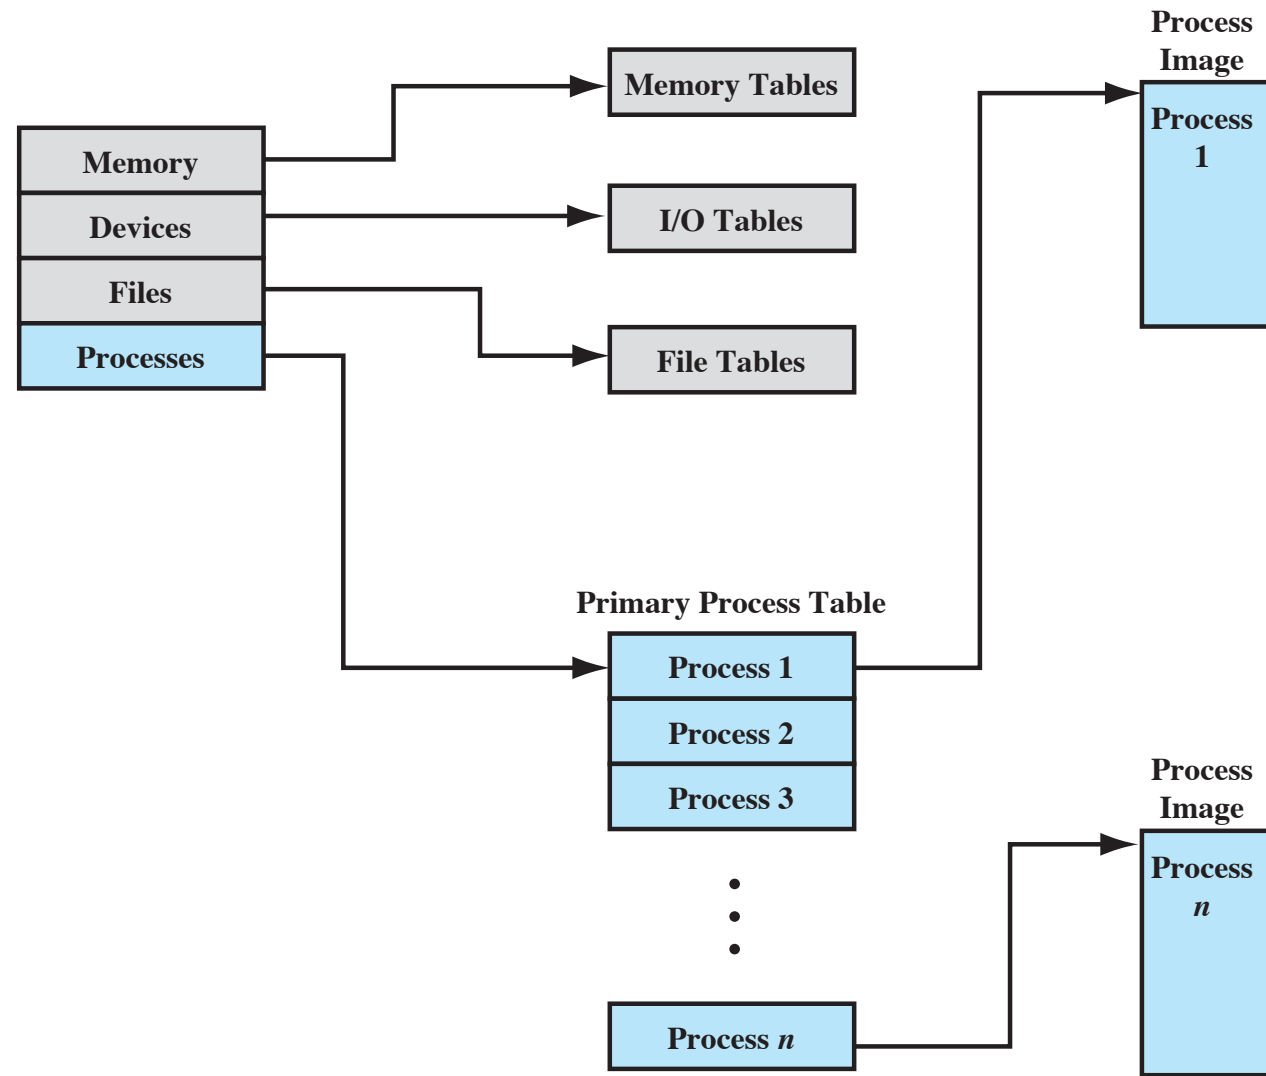
**Figure 3.7 Process States for Trace of Figure 3.4**





(b) With Two Suspend States

Figure 3.9 Process State Transition Diagram with Suspend States




**Figure 3.11 General Structure of Operating System Control Tables**

# Modes of Execution

## User Mode

- less-privileged mode
- user programs typically execute in this mode



## System Mode

- more-privileged mode
- also referred to as control mode or kernel mode
- kernel of the operating system



# Process Creation

- Once the OS decides to create a new process it:

assigns a unique process identifier to the new process

allocates space for the process

initializes the process control block

sets the appropriate linkages

creates or expands other data structures

# System Interrupts

## Interrupt

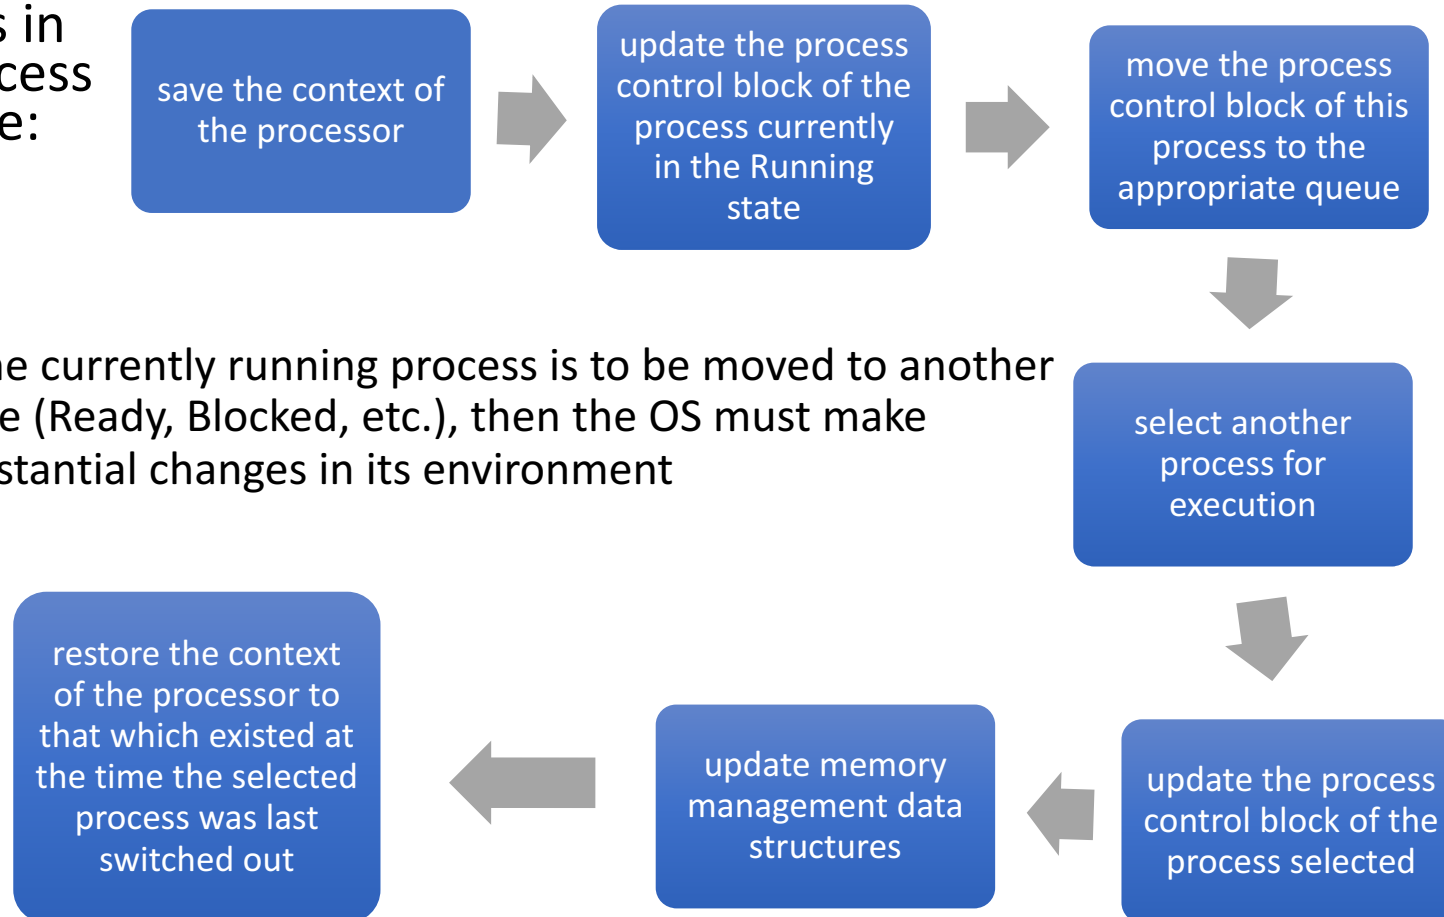
- Due to some sort of event that is external to and independent of the currently running process
  - clock interrupt
  - I/O interrupt
  - memory fault
- Time slice
  - the maximum amount of time that a process can execute before being interrupted

## Trap

- An error or exception condition generated within the currently running process
- OS determines if the condition is fatal
  - moved to the Exit state and a process switch occurs
  - action will depend on the nature of the error

# Change of Process State

- The steps in a full process switch are:



If the currently running process is to be moved to another state (Ready, Blocked, etc.), then the OS must make substantial changes in its environment

# Chapter 4

Threads and System Calls

# Processes and Threads

## Resource Ownership

Process includes a virtual address space to hold the process image

- the OS performs a protection function to prevent unwanted interference between processes with respect to resources

## Scheduling/Execution

Follows an execution path that may be interleaved with other processes


- a process has an execution state (Running, Ready, etc.) and a dispatching priority and is scheduled and dispatched by the OS



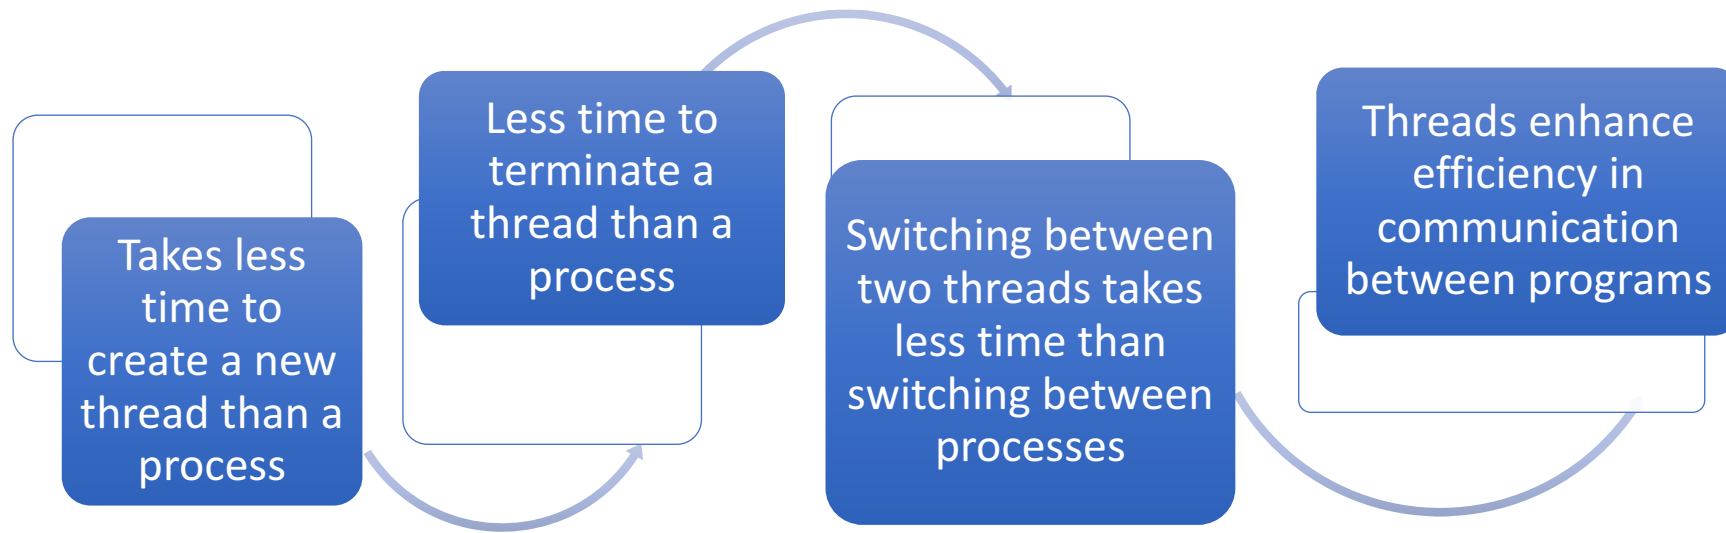
# One or More Threads in a Process

Each thread has:

- an execution state (Running, Ready, etc.)
- saved thread context when not running
- an execution stack
- some per-thread static storage for local variables
- access to the memory and resources of its process (all threads of a process share this)




# Benefits of Threads



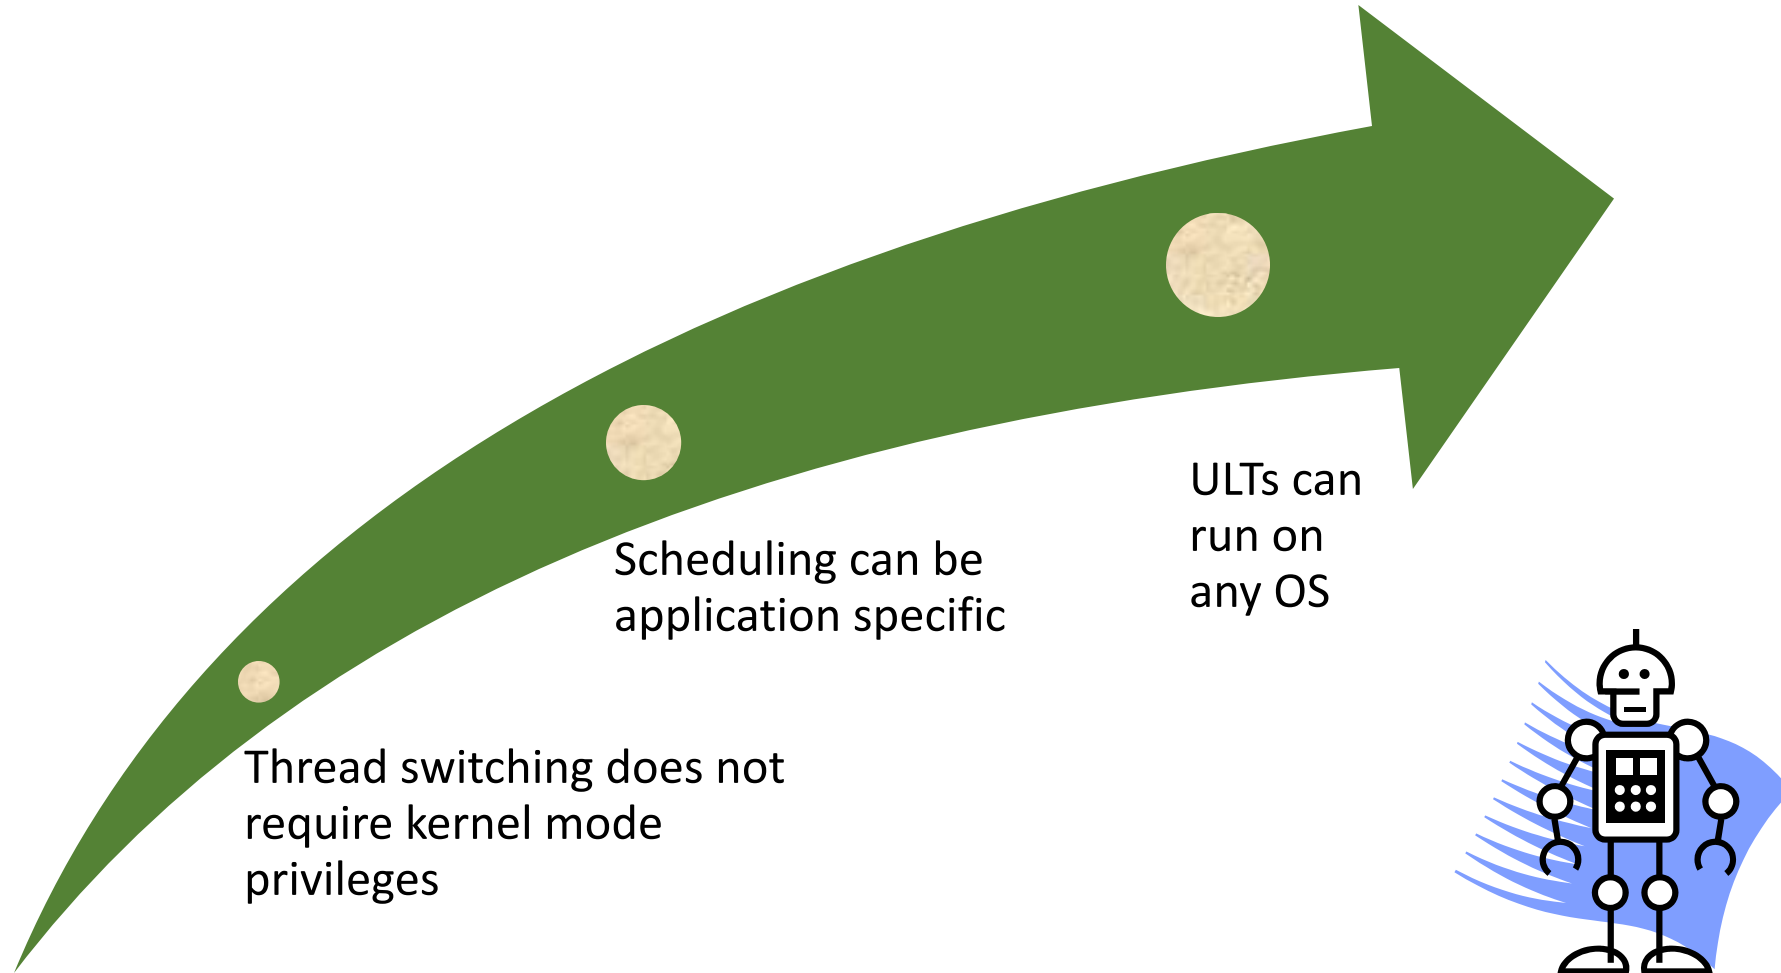
# User-Level Threads (ULTs)

- All thread management is done by the application
- The kernel is not aware of the existence of threads



(a) Pure user-level

# Advantages of ULTs



# Disadvantages of ULTs

- In a typical OS many system calls are blocking
  - as a result, when a ULT executes a system call, not only is that thread blocked, but all of the threads within the process are blocked
- In a pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing



# Overcoming ULT Disadvantages


## Jacketing

- converts a blocking system call into a non-blocking system call



Writing an application  
as multiple processes  
rather than multiple  
threads

# Kernel-Level Threads (KLTs)




(b) Pure kernel-level

- Thread management is done by the kernel
  - no thread management is done by the application
  - Windows is an example of this approach

# Advantages of KLTs

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines can be multithreaded





# Disadvantage of KLTs

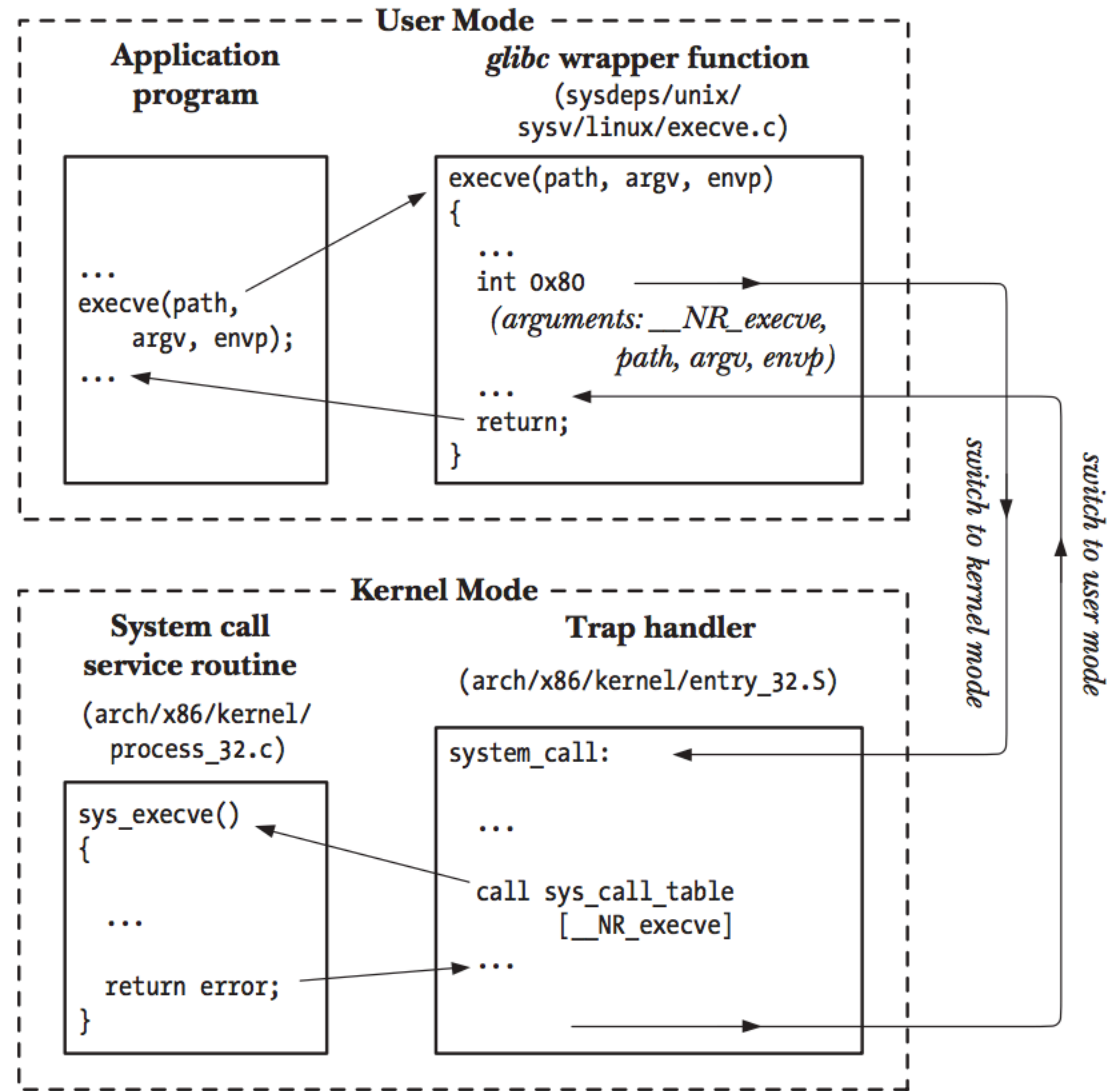
**\* The transfer of control from one thread to another within the same process requires a mode switch to the kernel**

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

**Table 4.1**  
**Thread and Process Operation Latencies ( $\mu$ s)**

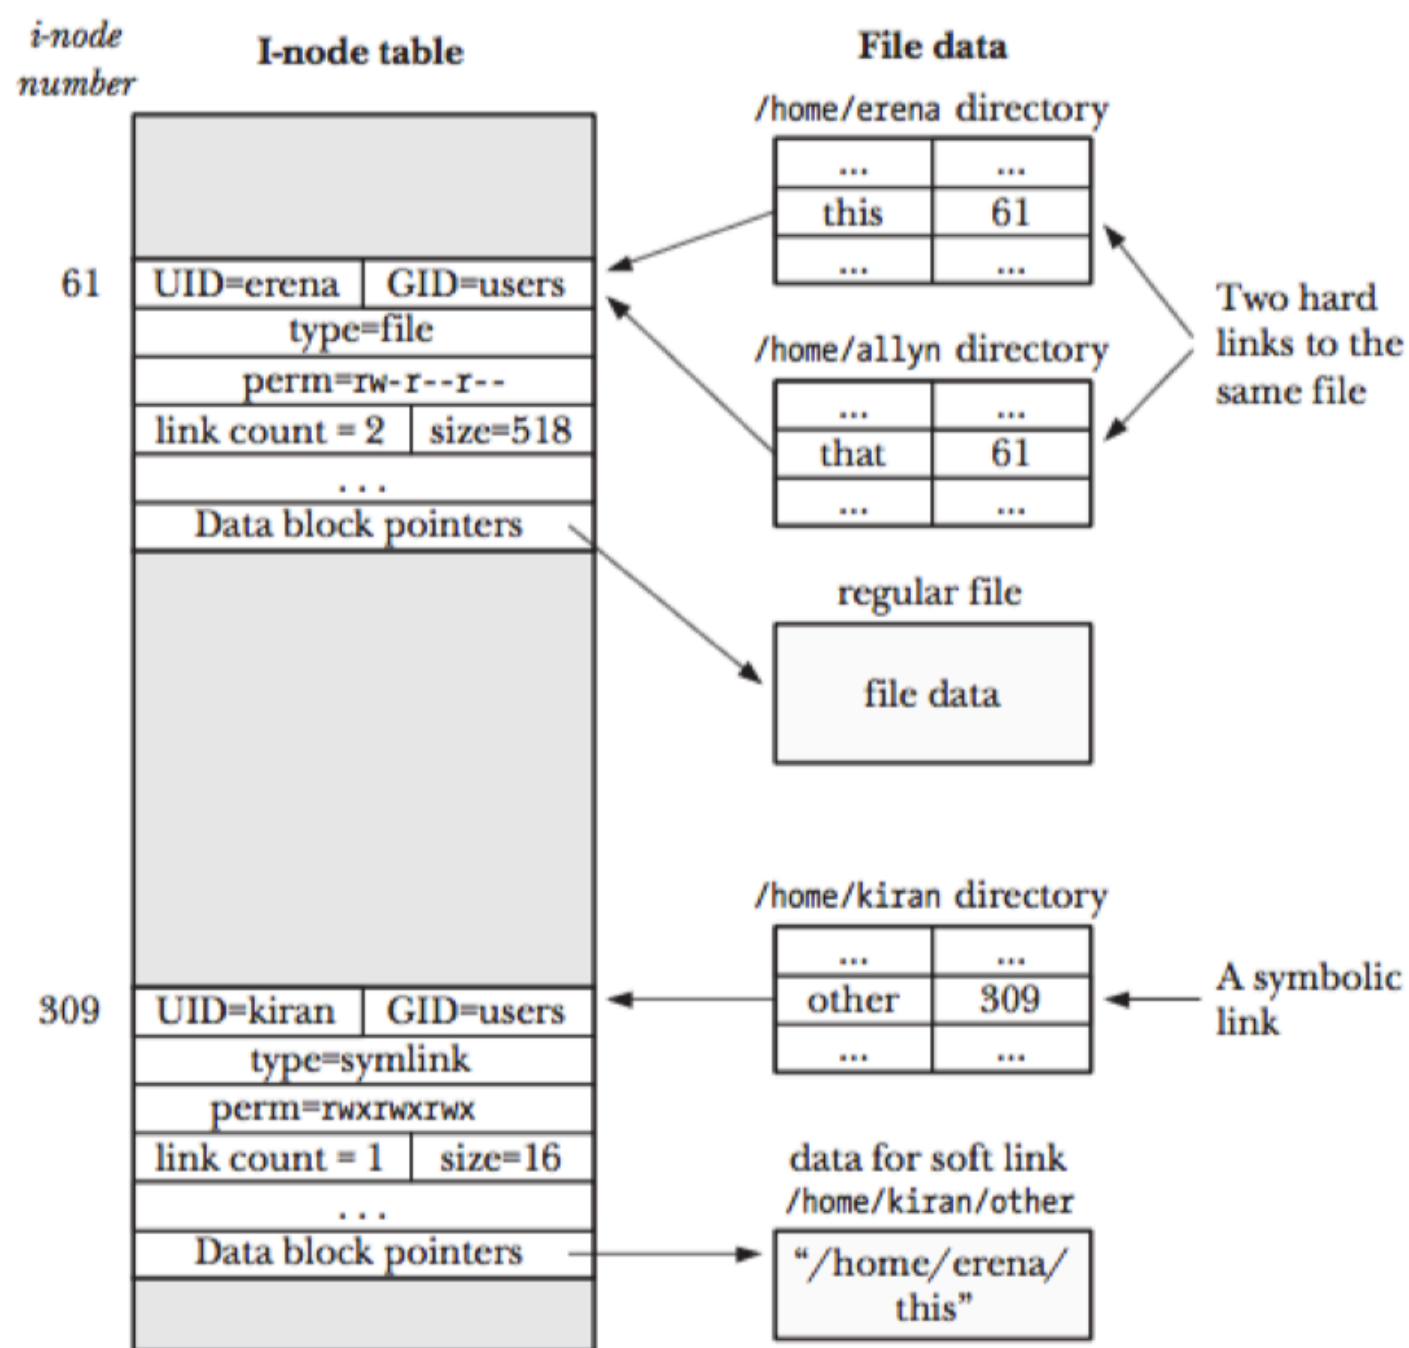
# Applications That Benefit

- Multithreaded native applications
  - characterized by having a small number of highly threaded processes
- Multiprocess applications
  - characterized by the presence of many single-threaded processes
- Java applications
- Multiinstance applications
  - multiple instances of the application in parallel



**Figure 3-1:** Steps in the execution of a system call





**Figure 18-2:** Representation of hard and symbolic links

# fork

```
#include <unistd.h>
```

```
pid_t fork(void);
```

In parent: returns process ID of child on success, or -1 on error;  
in successfully created child: always returns 0

```
pid_t childPid;          /* Used in parent after successful fork()
                           to record PID of child */
switch (childPid = fork()) {
case -1:                  /* fork() failed */
    /* Handle error */

case 0:                   /* Child of successful fork() comes here */
    /* Perform actions specific to child */

default:                  /* Parent comes here after successful fork() */
    /* Perform actions specific to parent */
}
```

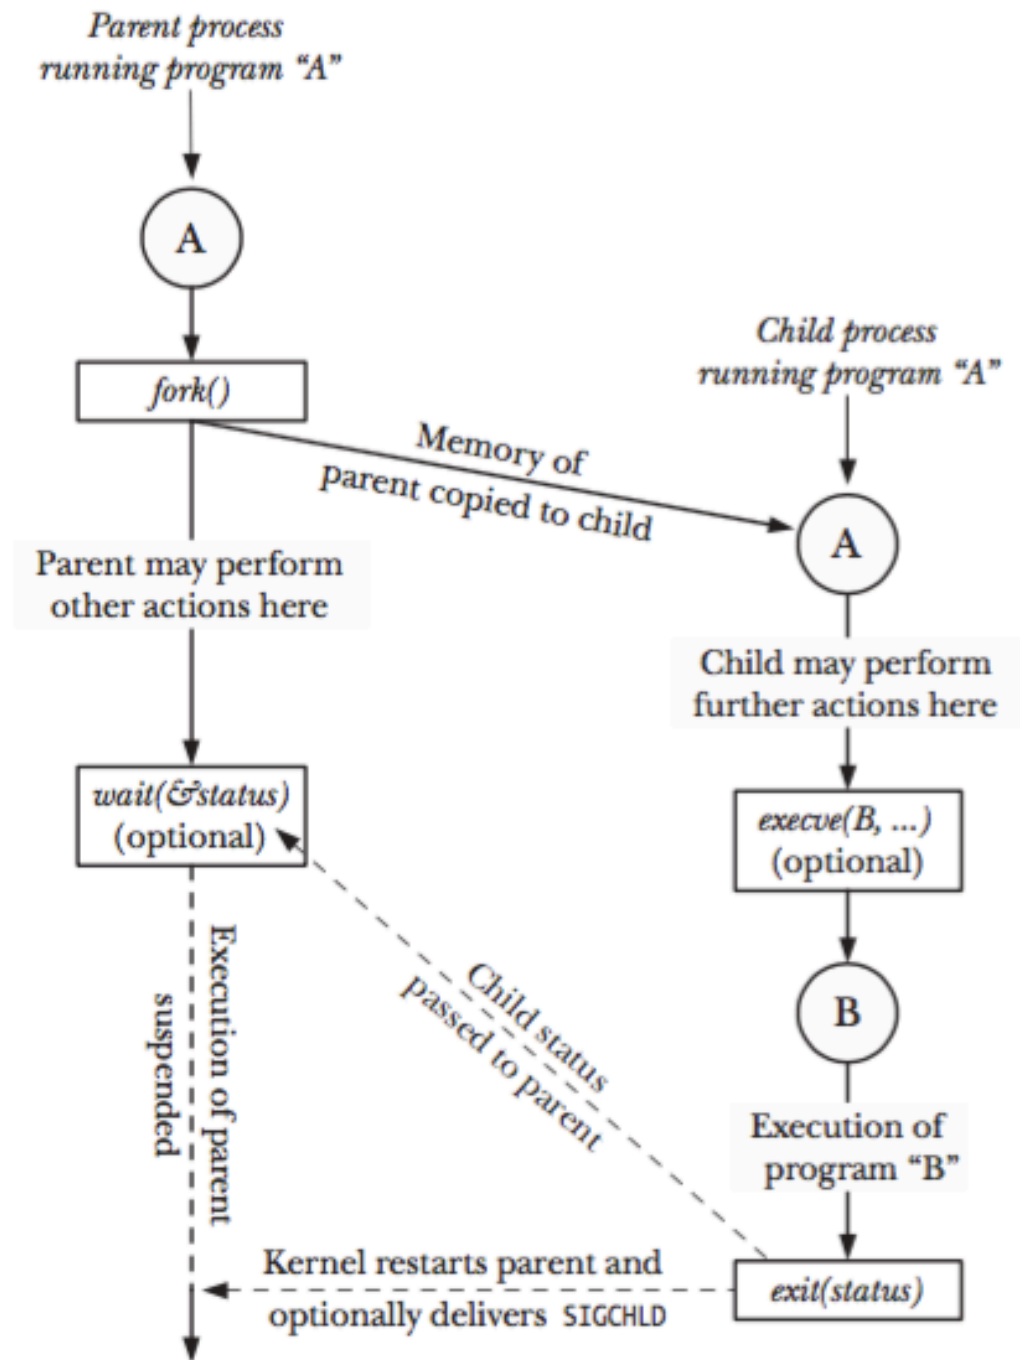
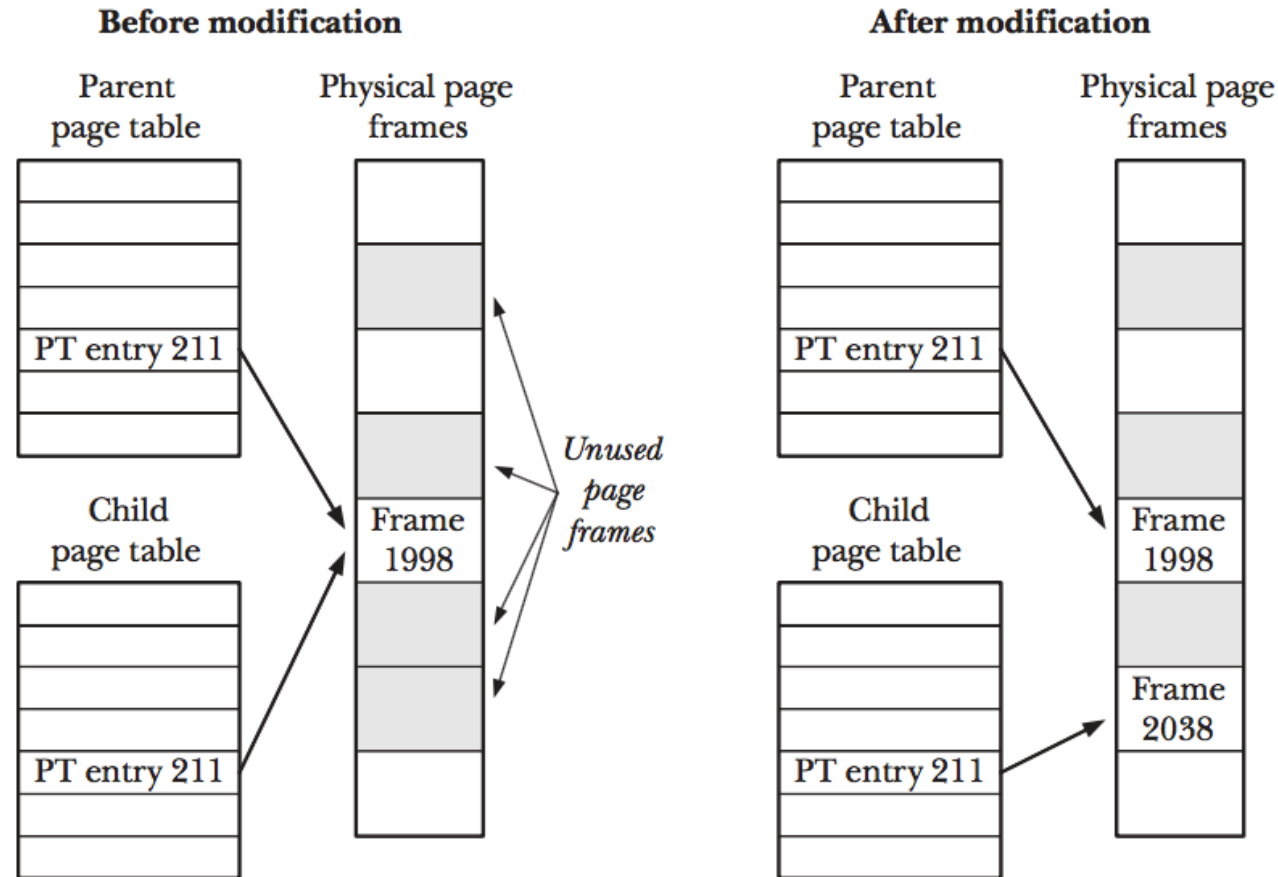


Figure 24-1: Overview of the use of *fork()*, *exit()*, *wait()*, and *execve()*

# Copy on write memory



**Figure 24-3:** Page tables before and after modification of a shared copy-on-write page




# Chapter 5

Concurrency

# Principles of Concurrency

- Interleaving and overlapping
  - can be viewed as examples of concurrent processing
  - both present the same problems
- Uniprocessor – the relative speed of execution of processes cannot be predicted
  - depends on activities of other processes
  - the way the OS handles interrupts
  - scheduling policies of the OS



# Difficulties of Concurrency

- Sharing of global resources
- Difficult for the OS to manage the allocation of resources optimally
- Difficult to locate programming errors as results are not deterministic and reproducible



# Race Condition

- Occurs when multiple processes or threads read and write data items
- The final result depends on the order of execution
  - the “loser” of the race is the process that updates last and will determine the final value of the variable



# Requirements for Mutual Exclusion

- Must be enforced
- A process that halts must do so without interfering with other processes
- No deadlock or starvation
- A process must not be denied access to a critical section when there is no other process using it
- No assumptions are made about relative process speeds or number of processes
- A process remains inside its critical section for a finite time only



# Mutual Exclusion: Hardware Support

## ■ Interrupt Disabling


- uniprocessor system
- disabling interrupts guarantees mutual exclusion

## ■ Disadvantages:

- the efficiency of execution could be noticeably degraded
- this approach will not work in a multiprocessor architecture

# Mutual Exclusion: Hardware Support

- Compare&Swap Instruction
  - also called a “compare and exchange instruction”
  - a **compare** is made between a memory value and a test value
  - if the values are the same a **swap** occurs
  - carried out atomically



# Semaphore

A variable that has an integer value upon which only three operations are defined:

- There is no way to inspect or manipulate semaphores other than these three operations

- 1) May be initialized to a nonnegative integer value
- 2) The semWait operation decrements the value
- 3) The semSignal operation increments the value



# Strong/Weak Semaphores

\* A queue is used to hold processes waiting on the semaphore

## ***Strong Semaphores***


- the process that has been blocked the longest is released from the queue first (FIFO)

## ***Weak Semaphores***

- the order in which processes are removed from the queue is not specified

# Monitor Characteristics

Local data variables are accessible only by the monitor's procedures and not by any external procedure

A dark gray arrow pointing downwards from the first box to the second box.

Process enters monitor by invoking one of its procedures

A dark blue arrow pointing downwards from the second box to the third box.

Only one process may be executing in the monitor at a time

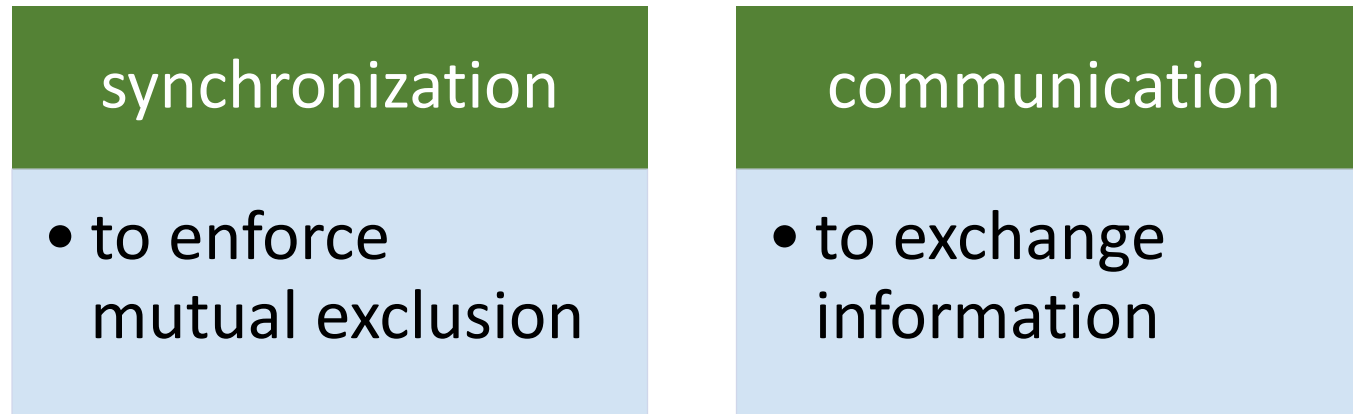
# Synchronization

- Achieved by the use of **condition variables** that are contained within the monitor and accessible only within the monitor
  - Condition variables are operated on by two functions:
    - cwait(c): suspend execution of the calling process on condition c
    - csignal(c): resume execution of some process blocked after a cwait on the same condition

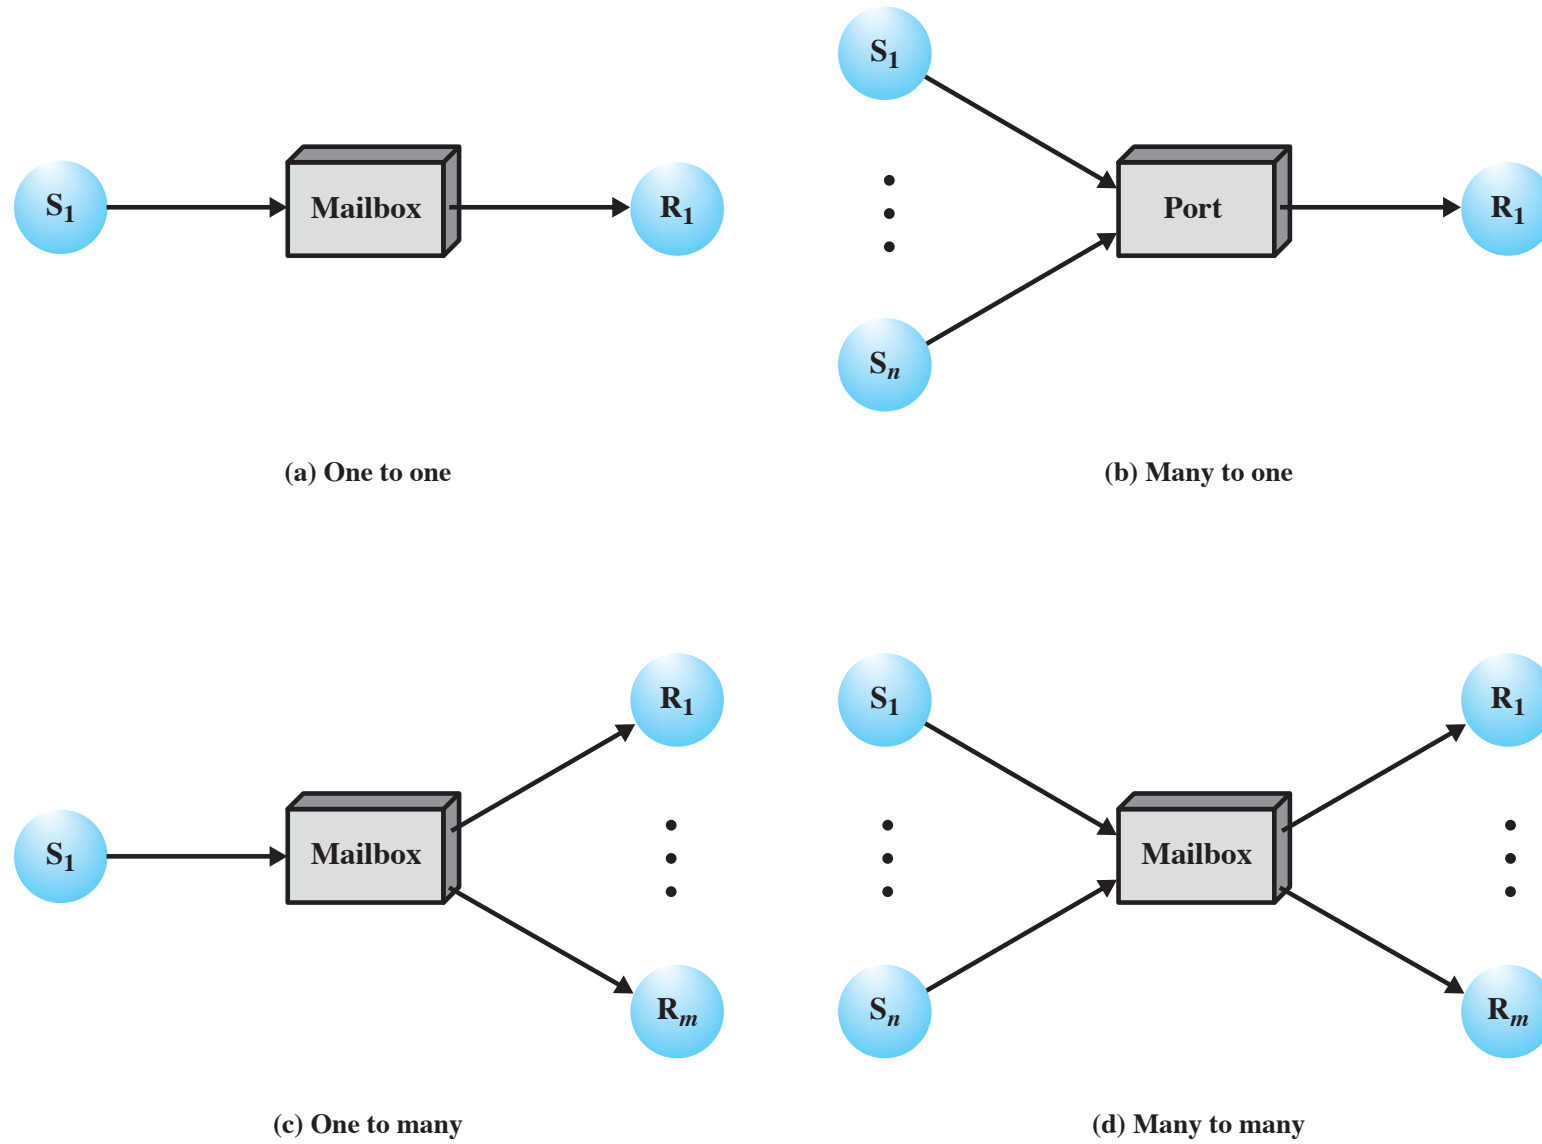


# Message Passing


- When processes interact with one another two fundamental requirements must be satisfied:




- Message Passing is one approach to providing both of these functions
  - works with distributed systems *and* shared memory multiprocessor and uniprocessor systems



**Figure 5.18 Indirect Process Communication**



**Figure 5.18 Indirect Process Communication**



**Figure 5.19 General Message Format**

# Chapter 6


Deadlock



# Deadlock

- The permanent blocking of a set of processes that either compete for system resources or communicate with each other
- A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set
- Permanent
- No efficient solution






**Figure 6.2 Example of Deadlock**

# Resource Categories

## Reusable

- can be safely used by only one process at a time and is not depleted by that use
- processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases, and semaphores




## Consumable


- one that can be created (produced) and destroyed (consumed)
  - interrupts, signals, messages, and information
  - in I/O buffers




(a) Resource is requested



(b) Resource is held



(c) Circular wait



(d) No deadlock

**Figure 6.5 Examples of Resource Allocation Graphs**

# Conditions for Deadlock

Mutual Exclusion	Hold-and-Wait	No Pre-emption	Circular Wait
<ul style="list-style-type: none"><li>• only one process may use a resource at a time</li></ul>	<ul style="list-style-type: none"><li>• a process may hold allocated resources while awaiting assignment of others</li></ul>	<ul style="list-style-type: none"><li>• no resource can be forcibly removed from a process holding it</li></ul>	<ul style="list-style-type: none"><li>• a closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain</li></ul>

## Banker's Algo

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0

**C - A**

R1	R2	R3
9	3	6

Resource vector **R**

R1	R2	R3
0	1	1

Available vector **V**

(a) Initial state

Figure 6.7 Determination of a Safe State

# Deadlock Strategies

Deadlock prevention strategies are very conservative

- limit access to resources by imposing restrictions on processes

Deadlock detection strategies do the opposite

- resource requests are granted whenever possible

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request matrix  
Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation  
matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource vector

R1	R2	R3	R4	R5
0	0	0	0	1

Allocation vector

Figure 6.10 Example for Deadlock Detection



# UNIX Concurrency Mechanisms

- UNIX provides a variety of mechanisms for interprocessor communication and synchronization including:

Pipes

Messages

Shared  
memory

Semaphores


Signals

# Chapter 7

Memory Management

# Memory Management Requirements

- Memory management is intended to satisfy the following requirements:
  - Relocation
  - Protection
  - Sharing
  - Logical organization
  - Physical organization



Technique	Description	Strengths	Weaknesses
<b>Fixed Partitioning</b>	Main memory is divided into a number of static partitions at system generation time. A process may be loaded into a partition of equal or greater size.	Simple to implement; little operating system overhead.	Inefficient use of memory due to internal fragmentation; maximum number of active processes is fixed.
<b>Dynamic Partitioning</b>	Partitions are created dynamically, so that each process is loaded into a partition of exactly the same size as that process.	No internal fragmentation; more efficient use of main memory.	Inefficient use of processor due to the need for compaction to counter external fragmentation.
<b>Simple Paging</b>	Main memory is divided into a number of equal-size frames. Each process is divided into a number of equal-size pages of the same length as frames. A process is loaded by loading all of its pages into available, not necessarily contiguous, frames.	No external fragmentation.	A small amount of internal fragmentation.
<b>Simple Segmentation</b>	Each process is divided into a number of segments. A process is loaded by loading all of its segments into dynamic partitions that need not be contiguous.	No internal fragmentation; improved memory utilization and reduced overhead compared to dynamic partitioning.	External fragmentation.
<b>Virtual Memory Paging</b>	As with simple paging, except that it is not necessary to load all of the pages of a process. Nonresident pages that are needed are brought in later automatically.	No external fragmentation; higher degree of multiprogramming; large virtual address space.	Overhead of complex memory management.
<b>Virtual Memory Segmentation</b>	As with simple segmentation, except that it is not necessary to load all of the segments of a process. Nonresident segments that are needed are brought in later automatically.	No internal fragmentation, higher degree of multiprogramming; large virtual address space; protection and sharing support.	Overhead of complex memory management.

Table 7.2

# Memory Management Techniques

(Table is on page 315 in textbook)

# Placement Algorithms

## Best-fit

- chooses the block that is closest in size to the request

## First-fit

- begins to scan memory from the beginning and chooses the first available block that is large enough

## Next-fit

- begins to scan memory from the location of the last placement and chooses the next available block that is large enough

# Buddy System

- Comprised of fixed and dynamic partitioning schemes
- Space available for allocation is treated as a single block
- Memory blocks are available of size  $2^K$  words,  $L \leq K \leq U$ , where
  - $2^L$  = smallest size block that is allocated
  - $2^U$  = largest size block that is allocated; generally  $2^U$  is the size of the entire memory available for allocation




# Paging

- Partition memory into equal fixed-size chunks that are relatively small
- Process is also divided into small fixed-size chunks of the same size

Pages	Frames
<ul style="list-style-type: none"><li>• chunks of a process</li></ul>	<ul style="list-style-type: none"><li>• available chunks of memory</li></ul>

# Page Table

- Maintained by operating system for each process
- Contains the frame location for each page in the process
- Processor must know how to access for the current process
- Used by processor to produce a physical address





# Segmentation

- A program can be subdivided into segments
  - may vary in length
  - there is a maximum length
- Addressing consists of two parts:
  - segment number
  - an offset
- Similar to dynamic partitioning
- Eliminates internal fragmentation

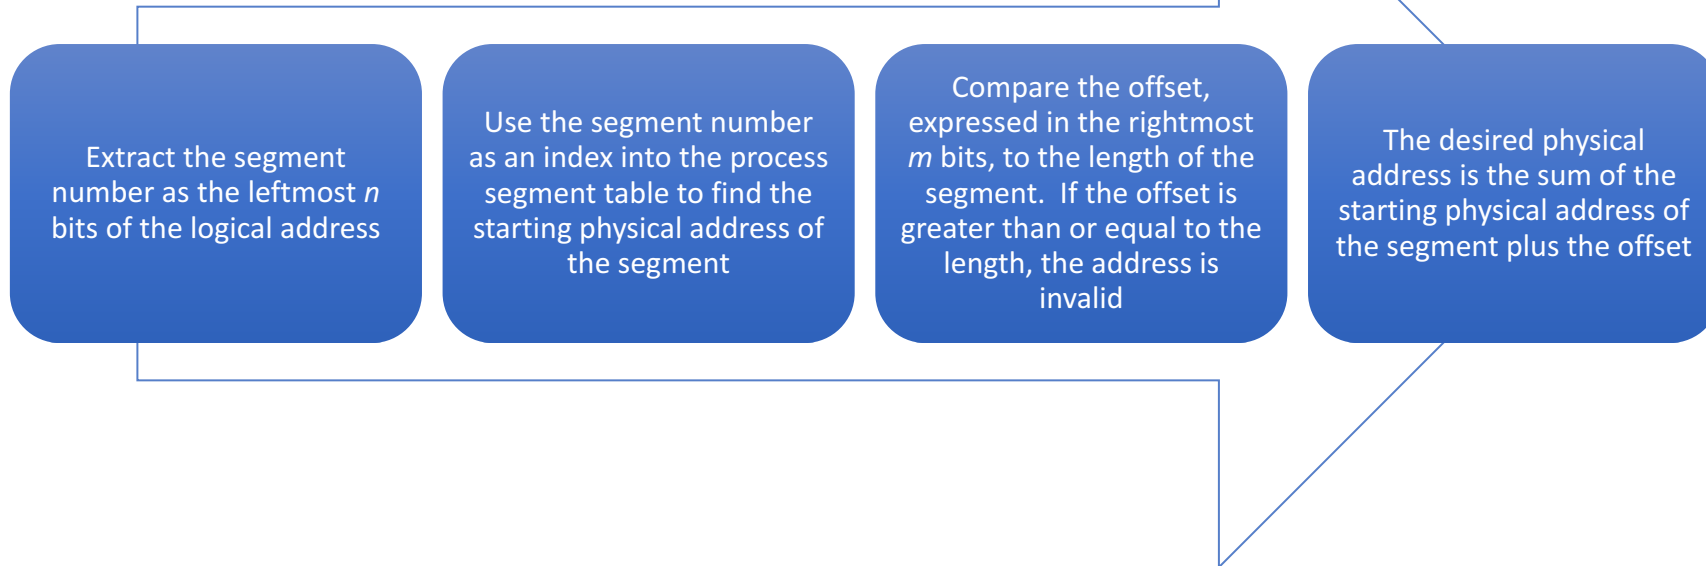


# Segmentation

- Usually visible
- Provided as a convenience for organizing programs and data
- Typically the programmer will assign programs and data to different segments
- For purposes of modular programming the program or data may be further broken down into multiple segments
  - the principal inconvenience of this service is that the programmer must be aware of the maximum segment size limitation

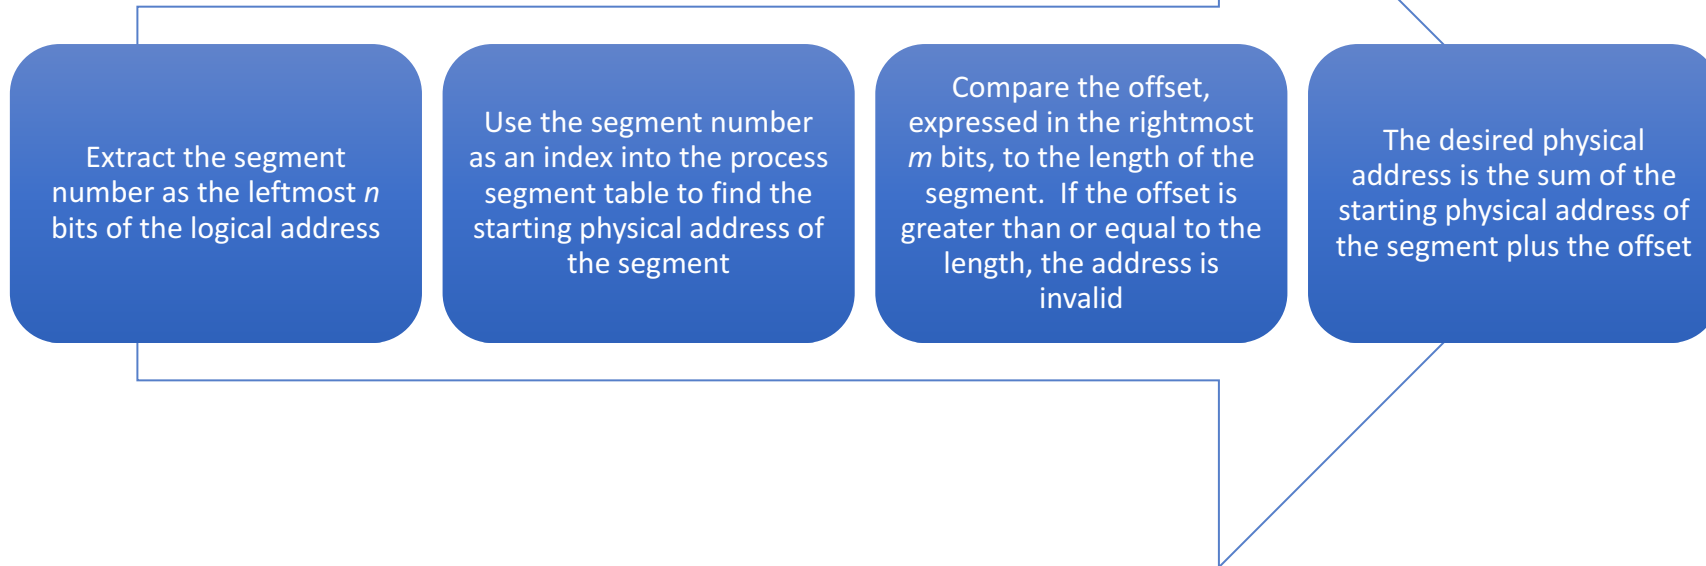
# Address Translation

- Another consequence of unequal size segments is that there is no simple relationship between logical addresses and physical addresses
- The following steps are needed for address translation:



# Address Translation

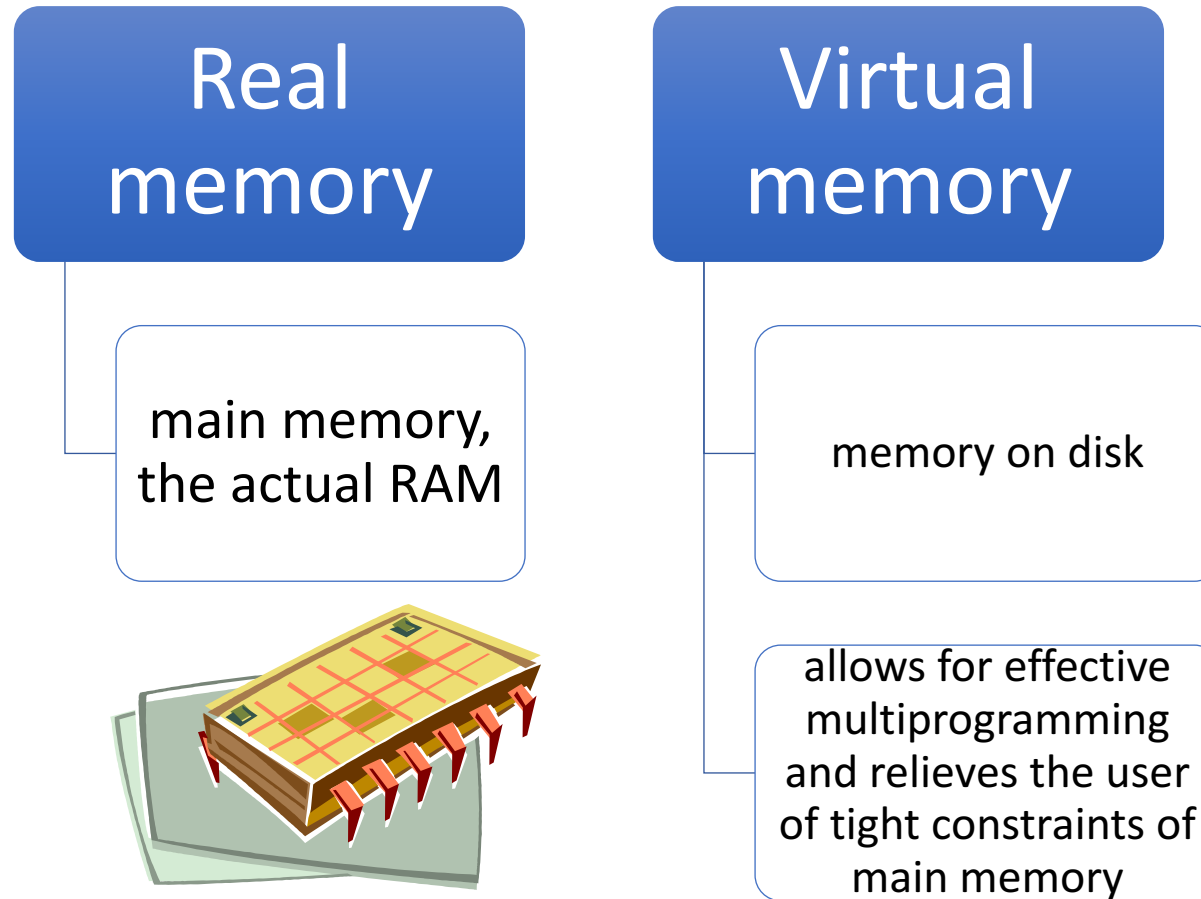
- Another consequence of unequal size segments is that there is no simple relationship between logical addresses and physical addresses
- The following steps are needed for address translation:



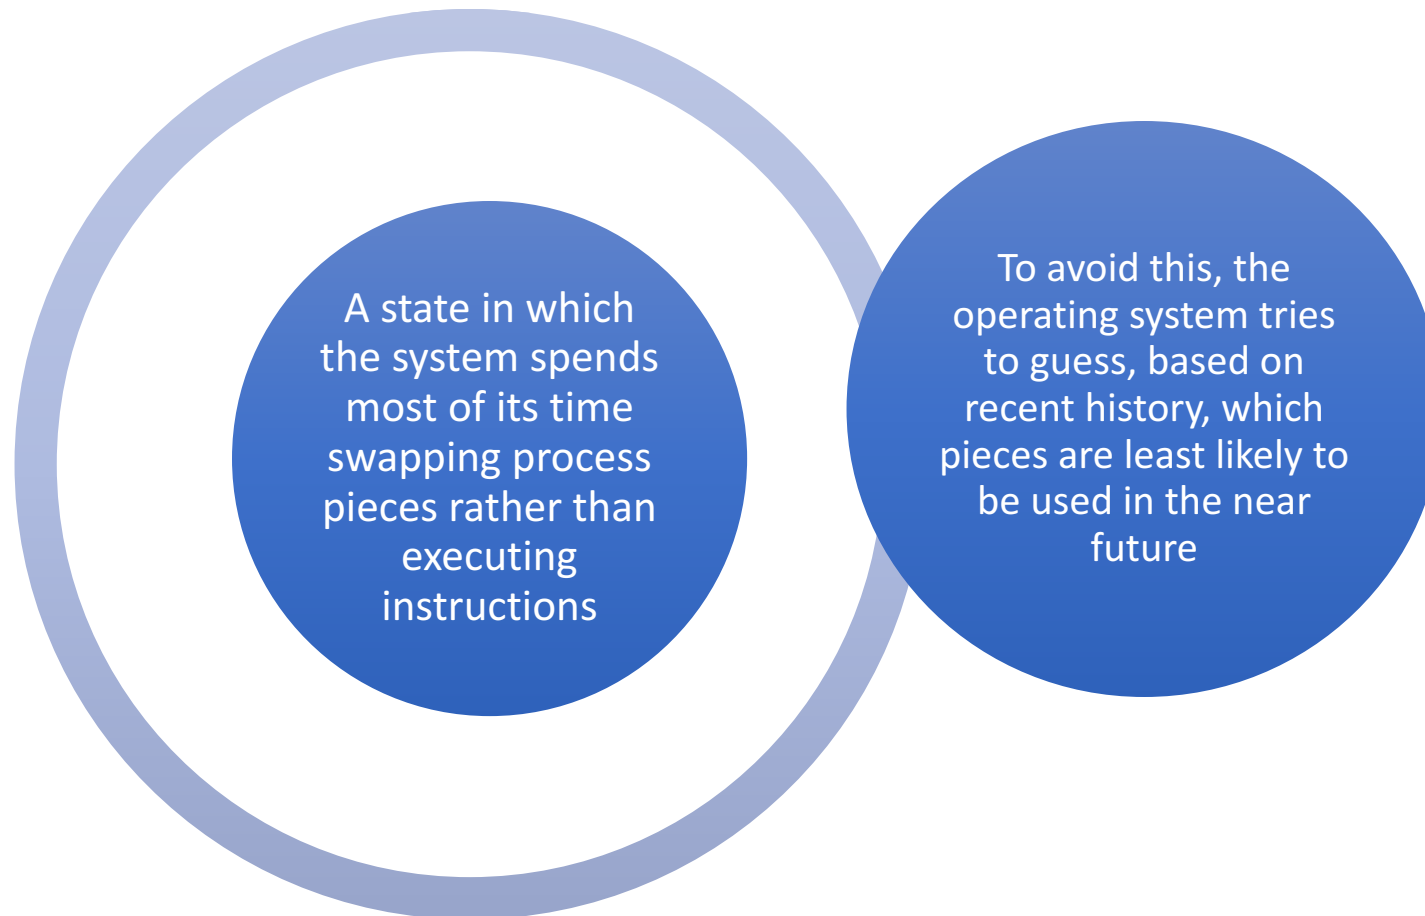
# Chapter 8

Virtual Memory and Virtual Page Allocation

# Real and Virtual Memory



# Thrashing



# Inverted Page Table

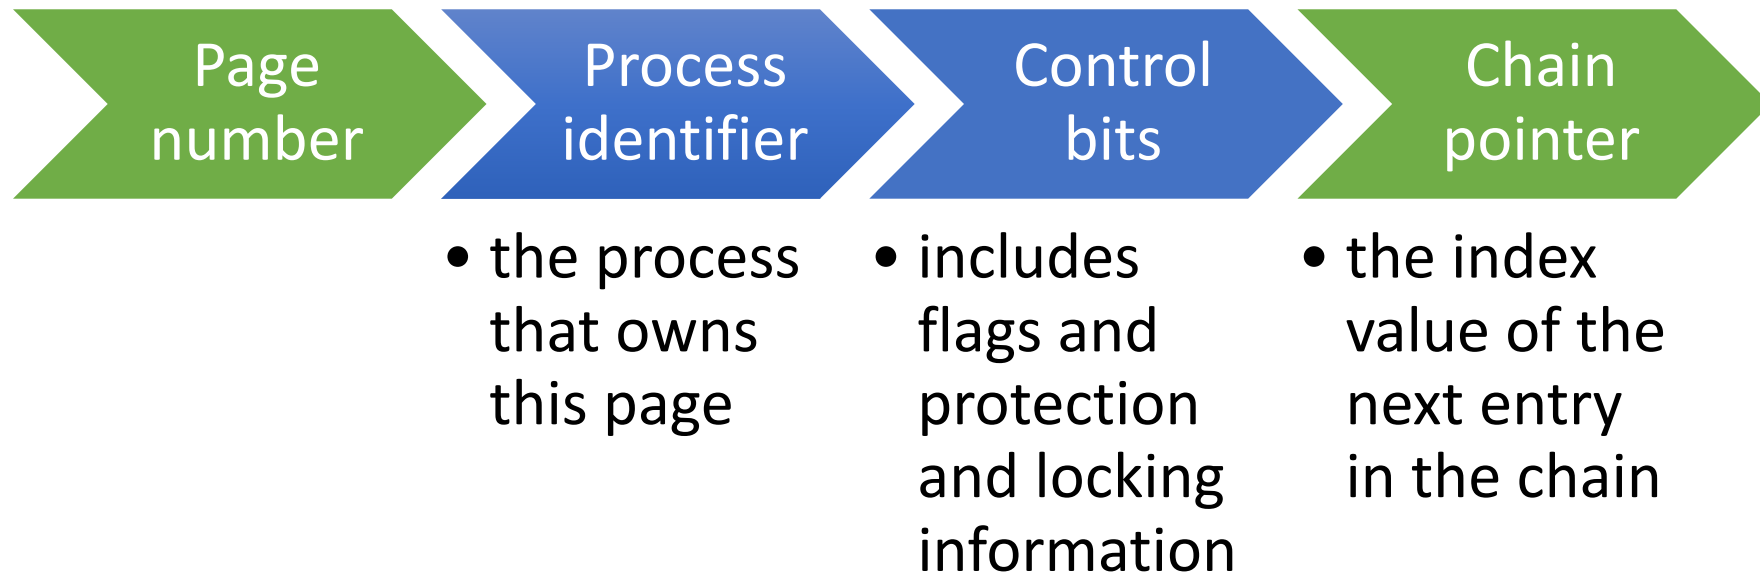
- Page number portion of a virtual address is mapped into a hash value
  - hash value points to inverted page table
- Fixed proportion of real memory is required for the tables regardless of the number of processes or virtual pages supported
- Structure is called *inverted* because it indexes page table entries by frame number rather than by virtual page number






# Inverted Page Table

Each entry in the page table includes:



# Translation Lookaside Buffer (TLB)

- Each virtual memory reference can cause two physical memory accesses:
  - one to fetch the page table entry
  - one to fetch the data
- To overcome the effect of doubling the memory access time, most virtual memory schemes make use of a special high-speed cache called a ***translation lookaside buffer***



F = page fault occurring after the frame allocation is initially filled

**Figure 8.14 Behavior of Four Page-Replacement Algorithms**